

# Log Instrumentation Specifications and Low Overhead Logging

Roy Shea

Department of Computer Science  
University of California, Los Angeles

`roy@cs.ucla.edu`

July 17, 2009

With Mani Srivastava, Young Cho, and endless input from folks in NESL.



# Observing the Unobservable

## The Vision

In-network algorithms directing a sensor network made up of thousands (or more) of sub-dollar wireless sensor nodes to localize a lion roaming through the jungle.

# Observing the Unobservable

## The Vision

In-network algorithms directing a sensor network made up of thousands (or more) of sub-dollar wireless sensor nodes to localize a lion roaming through the jungle.

## Problems with the Vision

# Observing the Unobservable

## The Vision

In-network algorithms directing a sensor network made up of thousands (or more) of sub-dollar wireless sensor nodes to localize a lion roaming through the jungle.

## Problems with the Vision

- Lions live in the savanna

# Observing the Unobservable

## The Vision

In-network algorithms directing a sensor network made up of thousands (or more) of sub-dollar wireless sensor nodes to localize a lion roaming through the jungle.

## Problems with the Vision

- Lions live in the savanna
- Still paying \$139.00 for a sensor node (TelosB with sensor suite)

# Observing the Unobservable

## The Vision

In-network algorithms directing a sensor network made up of thousands (or more) of sub-dollar wireless sensor nodes to localize a lion roaming through the jungle.

## Problems with the Vision

- Lions live in the savanna
- Still paying \$139.00 for a sensor node (TelosB with sensor suite)
- “Thousands” is a big number

# Observing the Unobservable

## The Vision

In-network algorithms directing a sensor network made up of thousands (or more) of sub-dollar wireless sensor nodes to localize a lion roaming through the jungle.

## Problems with the Vision

- Lions live in the savanna
- Still paying \$139.00 for a sensor node (TelosB with sensor suite)
- “Thousands” is a big number
- Easy to create suboptimal, fragile, and buggy localized algorithms

# After the Dream is Gone

## Picking Up Pieces



# After the Dream is Gone

## Picking Up Pieces

- Lions live in the savanna
  - ▶ Embrace breadth of expertise while seeking out researchers in other fields

# After the Dream is Gone

## Picking Up Pieces

- Lions live in the savanna
  - ▶ Embrace breadth of expertise while seeking out researchers in other fields
- Still paying \$139.00 for a sensor node (TelosB with sensor suite)
  - ▶ Count on the power of economy of scale *if* the demand is present

# After the Dream is Gone

## Picking Up Pieces

- Lions live in the savanna
  - ▶ Embrace breadth of expertise while seeking out researchers in other fields
- Still paying \$139.00 for a sensor node (TelosB with sensor suite)
  - ▶ Count on the power of economy of scale *if* the demand is present
- “Thousands” is a big number
  - ▶ Ends up that you can learn a lot about the world with one node, or two nodes, or ten nodes. . .

# After the Dream is Gone

## Picking Up Pieces

- Lions live in the savanna
  - ▶ Embrace breadth of expertise while seeking out researchers in other fields
- Still paying \$139.00 for a sensor node (TelosB with sensor suite)
  - ▶ Count on the power of economy of scale *if* the demand is present
- “Thousands” is a big number
  - ▶ Ends up that you can learn a lot about the world with one node, or two nodes, or ten nodes. . .
- Easy to create suboptimal, fragile, and buggy localized algorithms
  - ▶ Provide tools to help developers write efficient, robust, and functioning code for wireless embedded systems

# Observing the Unobservable

## Within the Observers of the Unobservable

### **My Vision**

Develop new techniques and tools that provide developers with the insight needed to create applications for and understand behavior within deployed networks of bottom tier sensing devices.

# Observing the Unobservable

## Within the Observers of the Unobservable

### My Vision

Develop new techniques and tools that provide developers with the insight needed to create applications for and understand behavior within deployed networks of bottom tier sensing devices.

#### LIS

Framework for describing and implementing logging tasks.

#### LowLog


Optimizing to token name spaces to reduce the overhead of collecting runtime call traces.

# Machine Class

## Desktop / Server



## Embedded Computer



## Microcontroller



### High End

- Native hardware support for debugging
- Ample resources that debugging infrastructure can use
- Rich I/O capabilities


### Bottom Tier

- Real time constraints limit interactive or CPU intensive debugging techniques
- Minimal resources require very small footprint debugging utilities


# Number of Machines (or Processes)



- Basic debugging setup
  - Ability to suspend program without impacting other processes
- 



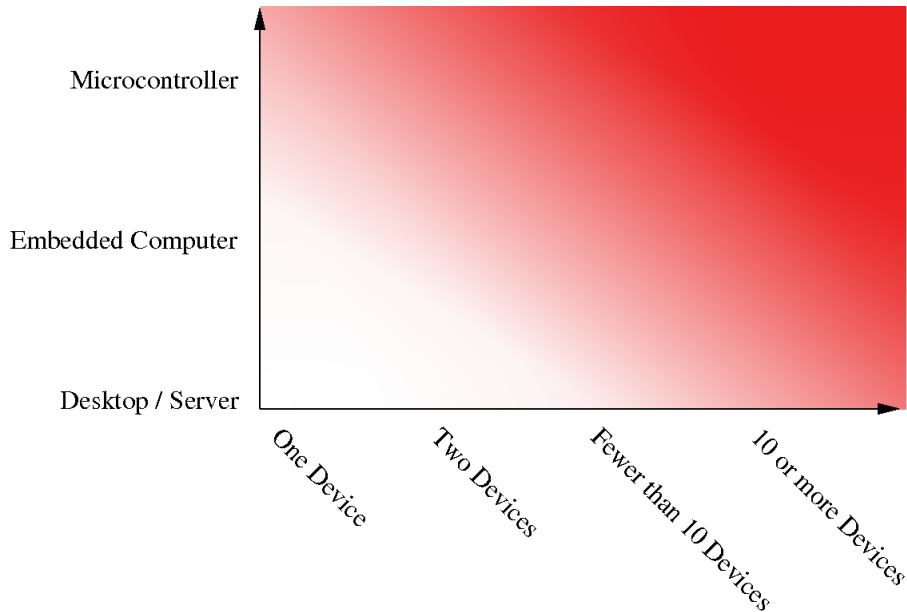
- Suspension of multiple processes tricky
  - Interactive debugging tracks execution on multiple processes to understand interactions
  - Logs capturing consistent system state are a powerful debugging aid
- 



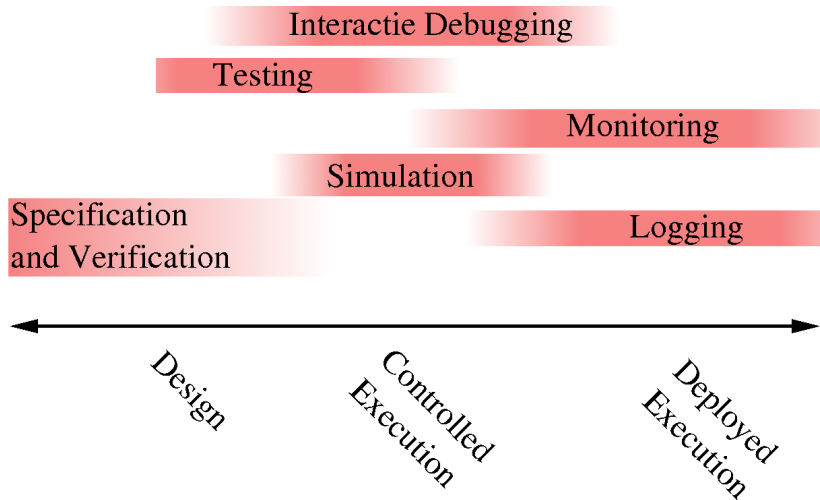
- Suspending execution becomes very difficult
- Challenging to capture consistent system state
- Use detailed logs used to understand aggregate system behavior



# Debugging Complexity of Different Systems



# Handling Bugs Through the Entire System Design Cycle



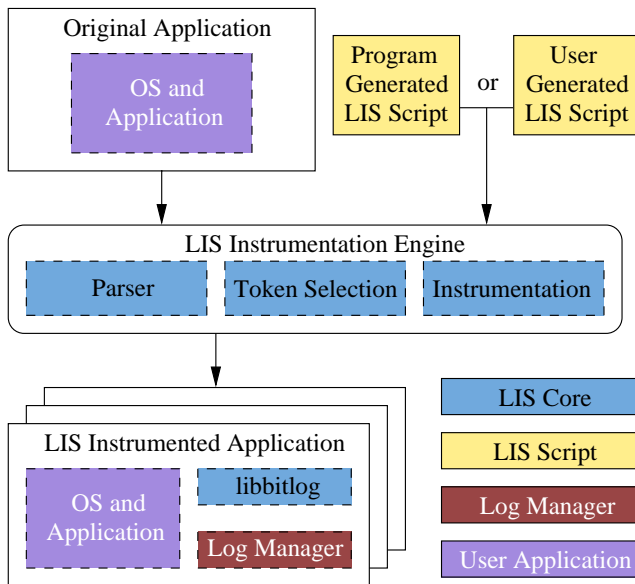
# Outline

- 1 Introduction
- 2 Log Instrumentation Specifications**
- 3 Low Overhead Logging with LIS
- 4 Conclusions

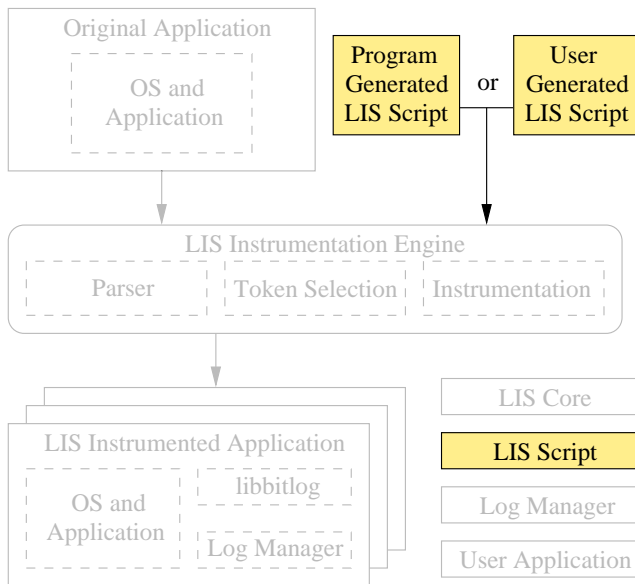
# Logging with Log Instrumentation Specifications (LIS)

- Observe demand for `printf` style logging in distributed embedded systems
  - ▶ Low learning curve and direct semantics make it easy to use
  - ▶ But resulting logs are typically verbose and ad-hoc logging is difficult to maintain
- Want to provide developers an alternate logging solution
  - ▶ Separation of logging specification from underlying code base
  - ▶ Encourages design of optimized logging tasks
  - ▶ Easy to pickup and use
- Logging streams of tokens via LIS is our answer

# Architecture



# Architecture



# Structure of a LIS Statement

## Statement Types

- Function header
- Function footer
- Function call
- Control flow
- Variable watchpoint

## Statement Location

- Union of log type and a function name
- Type specific directives bind log to specific called function, control flow type, or variable name

# Structure of a LIS Statement

## Statement Types

- Function header
- Function footer
- Function call
- Control flow
- Variable watchpoint

## Statement Location

- Union of log type and a function name
- Type specific directives bind log to specific called function, control flow type, or variable name

## Token Scope

**Global** Token value unique throughout the program

**Local** Token value is unique throughout the function

**Point** Token value is a singled fixed value



# LIS Specification

$Start \rightarrow Statements \mid \epsilon$

$Statements \rightarrow Stmt \ Statements \mid Stmt$

$Stmt \rightarrow Header \mid Footer \mid Call \mid ControlFlow \mid Watch$

$Header \rightarrow \mathbf{header} \ Placement \ Scope$

$Footer \rightarrow \mathbf{footer} \ Placement \ Scope$

$Call \rightarrow \mathbf{call} \ Placement \ Scope \ Target$

$ControlFlow \rightarrow \mathbf{controlflow} \ Placement \ Scope \ Flag \ Var$

$Watch \rightarrow \mathbf{watch} \ Placement \ Scope \ Var$

$Placement \rightarrow F$

$Scope \rightarrow \mathbf{global} \mid \mathbf{local} \mid \mathbf{point}$

$Target \rightarrow F \mid \mathbf{\_PTR\_}$

$Flag \rightarrow \mathbf{if} \mid \mathbf{switch} \mid \mathbf{loop} \mid \mathbf{if - switch} \mid \mathbf{if - loop}$   
 $\quad \mid \mathbf{switch - loop} \mid \mathbf{if - switch - loop}$

$Var \rightarrow \langle \text{Variable name from program} \rangle \mid \mathbf{\_ANY\_}$

$F \rightarrow \langle \text{Function name from program} \rangle$

# LIS Script

## Script and Target Program

```
header read_done global
controlflow read_done local if send_busy
footer read_done point
```

```
/* Pre-LIS */
void read_done(error_t result, uint16_t data) {

    if (send_busy == TRUE) {

        return;
    }

    /* Rest of function body elided... */

    return;
}
```

# LIS Script

## Instrumented Program

```
header read_done global
controlflow read_done local if send_busy
footer read_done point
```

```
/* Post-LIS */
void read_done(error_t result, uint16_t data) {

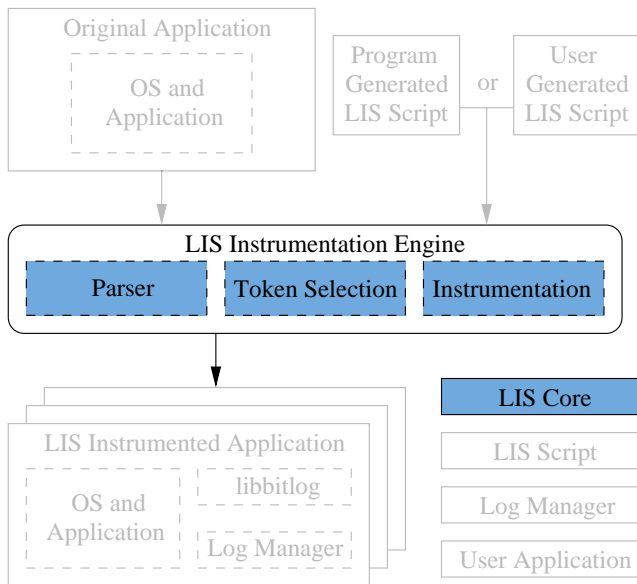
    bitlog_write(4, 3); /* Header LIS statement */

    if (send_busy == TRUE) {
        bitlog_write(5, 3); /* Control flow LIS statement */
        bitlog_write(0, 1); /* Footer LIS statement */
        return;
    }
    bitlog_write(6, 3); /* Control flow LIS statement */

    /* Rest of function body elided... */

    bitlog_write(0, 1); /* Footer LIS statement */
    return;
}
```

# Architecture



# Instantiation of LIS

## Instrumentation Engine

- Built using using C Intermediary Language (CIL) framework
  - ▶ Input a LIS script and source program
  - ▶ Outputs instrumented program
- Functional ports for x86, ATmega128, and MSP430 chip sets
- Patch available to integrate LIS into the TinyOS build system

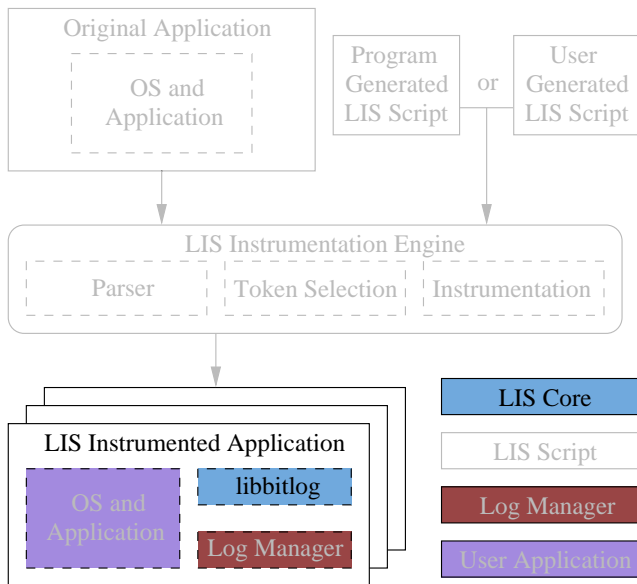
# Instantiation of LIS

## Instrumentation Engine

- Built using using C Intermediary Language (CIL) framework
  - ▶ Input a LIS script and source program
  - ▶ Outputs instrumented program
- Functional ports for x86, ATmega128, and MSP430 chip sets
- Patch available to integrate LIS into the TinyOS build system

|                             | Mica2          |      | MicaZ          |      | TelosB         |      |
|-----------------------------|----------------|------|----------------|------|----------------|------|
|                             | Program Memory | RAM  | Program Memory | RAM  | Program Memory | RAM  |
| <b>TinyOS Radio Stack</b>   | 7178           | 137  | 9264           | 210  | 8456           | 229  |
| <b>TinyOS CTP</b>           | 9692           | 1230 | 10284          | 1360 | 11126          | 1295 |
| <b>LogTap (CTP)</b>         | 1384           | 303  | 1412           | 351  | 2228           | 353  |
| <b>LogTap (broadcast)</b>   | 108            | 113  | 74             | 128  | 388            | 131  |
| <b>Bitlog Library</b>       | 386            | 43   | 386            | 43   | 362            | 43   |
| <b>Call to bitlog_write</b> | 14             | 0    | 14             | 0    | 12             | 0    |

# Architecture



# Instantiation of LIS

## Runtime Support

### Bitlog Library

- Writing of tokens into a buffer managed by a logging library
- Includes the Bitlog library that provides low overhead bit aligned logging

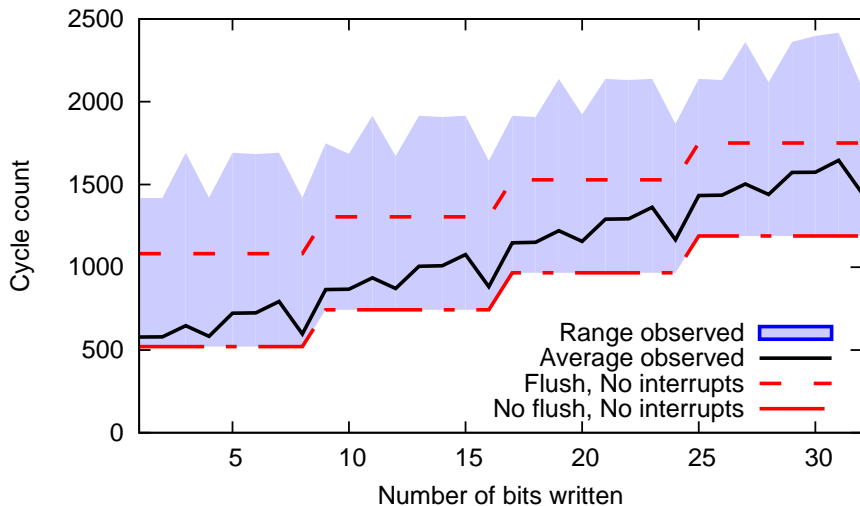
### LogTap

- Storage / transfer of logs managed by log management library
- Writing log buffers out to stderr for use on desktop machines
- Broadcasting logs using the TinyOS AMSend interface
- Routing logs using the TinyOS collection tree protocol (CTP)



# Instantiation of LIS

## Runtime Support

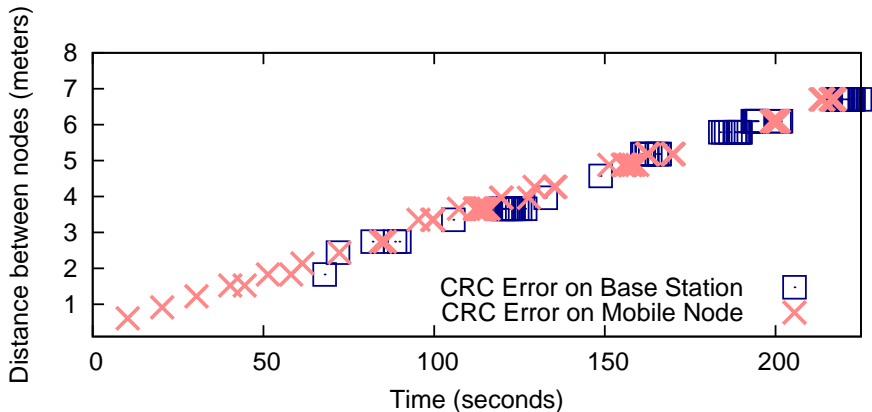


# Using LIS: Observing CRC Errors

```
controlflow CC2420ReceiveP.RXFIFO.readDone global if buf
```

# Using LIS: Observing CRC Errors

```
controlflow CC2420ReceiveP.RXFIFO.readDone global if buf
```



## Using LIS: Learning about CTP

```
header CtpForwardingEngineP.0.sendTask.runTask point  
controlflow CtpForwardingEngineP.0.sendTask.runTask \  
    local if-switch-loop __ANY__
```

## Using LIS: Learning about CTP

```
header CtpForwardingEngineP.0.sendTask.runTask point  
controlflow CtpForwardingEngineP.0.sendTask.runTask \  
    local if-switch-loop __ANY__
```

| LIS Log Token | Comments                                   |
|---------------|--|
| <--           | Enter runTask                              |
| BID: 1        | CTP is not busy                            |
| BID: 3        | Send queue is nonempty                     |
| BID: 32       | Node is not root so must route message     |
| BID: 33       | Found a route to the root                  |
| BID: 4        | Prepare to send data                       |
| BID: 9        | Neighbor is not congested                  |
| BID: 16       | Message has not already been sent          |
| BID: 19       | Node is not the root so must route message |
| BID: 21       | Found current path quality metric          |
| BID: 23       | Not congested                              |
| BID: 24       | Succeeded in sending message               |
| BID: 25       | Note that the client sent a packet         |

# Outline

- 1 Introduction
- 2 Log Instrumentation Specifications
- 3 Low Overhead Logging with LIS**
- 4 Conclusions

# Tracing Calls Within a Region of Interest with LowLog

- Function call traces
  - ▶ Describe the sequence of function calls made at runtime
  - ▶ Already used by developers to understand system behavior
  - ▶ Often provides enough information in itself to diagnose bugs or jump start more aggressive debugging efforts
- Region of interest (ROI)
  - ▶ Subsystem of interest to a developer
  - ▶ For example there is (usually) no reason to trace execution through the kernel when debugging a user space application
  - ▶ Provides a focused view of a system
- LowLog provides optimized region of interest call tracing
- LowLog sits as a higher level analysis that outputs efficient LIS scripts


# Partitioning a Program into an ROI

- *Entry function* is reachable from a function not within the ROI
- *Body function* only reachable from functions within the ROI
- *Return of interest* when any function from the ROI returns




# Partitioning a Program into an ROI

- *Entry function* is reachable from a function not within the ROI
- *Body function* only reachable from functions within the ROI
- *Return of interest* when any function from the ROI returns




# Partitioning a Program into an ROI

- *Entry function* is reachable from a function not within the ROI
- *Body function* only reachable from functions within the ROI
- *Return of interest* when any function from the ROI returns



## Detail



## Create Log of Entry, Body, and Return Tokens

# Optimizing Call Trace Bandwidth


- Standard call tracing uses *global identifier logging* that lumps entry and body tokens into a single name space and use simple fixed bit width token encoding
- LowLog proposes to alternate schemes
  - ▶ *Local identifier logging* creates multiple caller specific token name spaces
  - ▶ *Control flow logging* tracks runtime control flow decisions rather than body calls
- All three logging techniques can apply more powerful encoding techniques if more information is available

# Optimizing Control Flow Logging

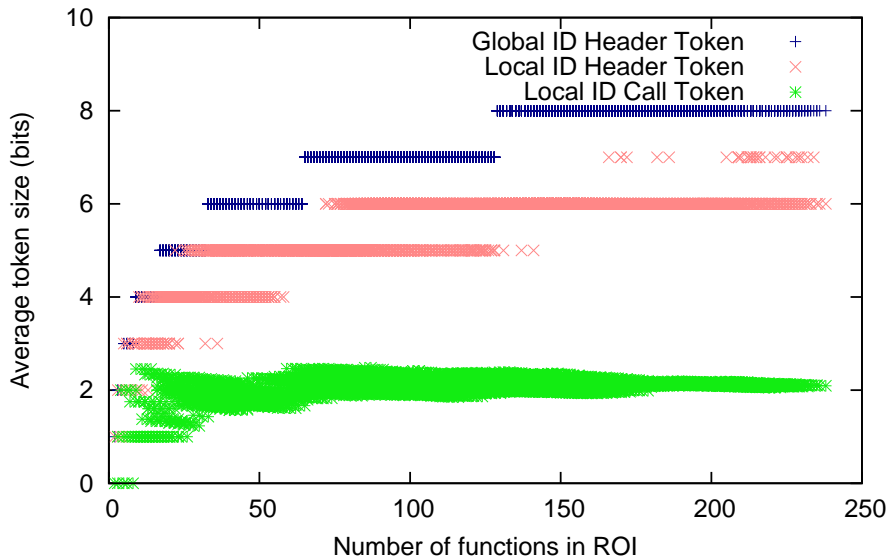
## Only Log Control Flow Decisions Effecting Called Functions

- Reverse dataflow analysis to track sets of functions that must be called after reaching a statement
- Join function examines sets of called functions
  - ▶ If the sets are different then logging of the control flow taken from the current node is tracked and the empty set is passed into the transfer function
  - ▶ If the sets are identical then the set (either since they are identical) is passed into the transfer function
- Transfer function prepends called functions from the current node to the list of called functions


# Call Trace Bandwidth



# Effect of ROI Size on Token Width



# Effect of ROI Size on Bandwidth





# Using LowLog: Problems in the Radio Stack

CC2420TransmitP

# Using LowLog: Problems in the Radio Stack

## CC2420TransmitP

- TinyOS runtime resource accounting and message time stamping each work in isolation
- Combining the two causes many time stamps to be lost
- Gathered call traces from the TinyOS CC2420TransmitP component
  - ▶ Traces with both resource accounting and time stamping reveal a frequent number of calls to `CC2420Receive.sfd_dropped`
  - ▶ Insight limited code search down to about 20 lines of code
- Timing delays introduced by resource accounting cause an overly conservative block of code to invalidate time stamps

# Outline

- 1 Introduction
- 2 Log Instrumentation Specifications
- 3 Low Overhead Logging with LIS
- 4 Conclusions**

# Summary

- LIS and LowLog help developers understand what their system is doing
  - ▶ LIS provides developers with a convenient and powerful framework for describing logging tasks
  - ▶ LowLog sits upon LIS and provides optimized call trace logging
- Have not yet found any lions, but am still looking

**Want to play with LIS and LowLog?**

`http://nesl.ee.ucla.edu/research/lis`

# Questions?